

# Assignment 4 - Naive Machine Translation and LSH

You will now implement your first machine translation system and then you will see how locality sensitive hashing works. Let's get started by importing the required functions!

If you are running this notebook in your local computer, don't forget to download the twitter samples and stopwords from nltk.

```
nltk.download('stopwords')
nltk.download('twitter_samples')
```

**NOTE:** The Exercise xx numbers in this assignment *are inconsistent* with the UNQ\_Cx numbers.

## This assignment covers the following topics:

- [1. The word embeddings data for English and French words](#)
  - [1.1 Generate embedding and transform matrices](#)
    - [Exercise 1](#)
- [2. Translations](#)
  - [2.1 Translation as linear transformation of embeddings](#)
    - [Exercise 2](#)
    - [Exercise 3](#)
    - [Exercise 4](#)
  - [2.2 Testing the translation](#)
    - [Exercise 5](#)
    - [Exercise 6](#)
- [3. LSH and document search](#)
  - [3.1 Getting the document embeddings](#)
    - [Exercise 7](#)
    - [Exercise 8](#)
  - [3.2 Looking up the tweets](#)
  - [3.3 Finding the most similar tweets with LSH](#)
  - [3.4 Getting the hash number for a vector](#)
    - [Exercise 9](#)
  - [3.5 Creating a hash table](#)
    - [Exercise 10](#)
  - [3.6 Creating all hash tables](#)
    - [Exercise 11](#)

```
In [1]: import pdb
import pickle
import string

import time

import gensim
import matplotlib.pyplot as plt
import nltk
import numpy as np
import scipy
import sklearn
from gensim.models import KeyedVectors
from nltk.corpus import stopwords, twitter_samples
from nltk.tokenize import TweetTokenizer

from utils import (cosine_similarity, get_dict,
                   process_tweet)
from os import getcwd
```

```
In [2]: # add folder, tmp2, from our Local workspace containing pre-downloaded corpora f
filePath = f"{getcwd()}/../tmp2/"
nltk.data.path.append(filePath)
```

## 1. The word embeddings data for English and French words

Write a program that translates English to French.

### The data

The full dataset for English embeddings is about 3.64 gigabytes, and the French embeddings are about 629 megabytes. To prevent the Coursera workspace from crashing, we've extracted a subset of the embeddings for the words that you'll use in this assignment.

If you want to run this on your local computer and use the full dataset, you can download the

- English embeddings from Google code archive word2vec [look for GoogleNews-vectors-negative300.bin.gz \(https://code.google.com/archive/p/word2vec/\)](https://code.google.com/archive/p/word2vec/)
  - You'll need to unzip the file first.
- and the French embeddings from [cross\\_lingual\\_text\\_classification \(https://github.com/vjstark/crosslingual\\_text\\_classification\)](https://github.com/vjstark/crosslingual_text_classification).
  - in the terminal, type (in one line) curl -o ./wiki.multi.fr.vec <https://dl.fbaipublicfiles.com/arrival/vectors/wiki.multi.fr.vec>

Then copy-paste the code below and run it.

```
# Use this code to download and process the full dataset on your Local c
```

*omputer*

```

from gensim.models import KeyedVectors

en_embeddings = KeyedVectors.load_word2vec_format('./GoogleNews-vectors-negative300.bin', binary = True)
fr_embeddings = KeyedVectors.load_word2vec_format('./wiki.multi.fr.vec')

# Loading the english to french dictionaries
en_fr_train = get_dict('en-fr.train.txt')
print('The length of the english to french training dictionary is', len(en_fr_train))
en_fr_test = get_dict('en-fr.test.txt')
print('The length of the english to french test dictionary is', len(en_fr_train))

english_set = set(en_embeddings.vocab)
french_set = set(fr_embeddings.vocab)
en_embeddings_subset = {}
fr_embeddings_subset = {}
french_words = set(en_fr_train.values())

for en_word in en_fr_train.keys():
    fr_word = en_fr_train[en_word]
    if fr_word in french_set and en_word in english_set:
        en_embeddings_subset[en_word] = en_embeddings[en_word]
        fr_embeddings_subset[fr_word] = fr_embeddings[fr_word]

for en_word in en_fr_test.keys():
    fr_word = en_fr_test[en_word]
    if fr_word in french_set and en_word in english_set:
        en_embeddings_subset[en_word] = en_embeddings[en_word]
        fr_embeddings_subset[fr_word] = fr_embeddings[fr_word]

pickle.dump( en_embeddings_subset, open( "en_embeddings.p", "wb" ) )
pickle.dump( fr_embeddings_subset, open( "fr_embeddings.p", "wb" ) )

```

## The subset of data

To do the assignment on the Coursera workspace, we'll use the subset of word embeddings.

In [3]:

```
en_embeddings_subset = pickle.load(open("en_embeddings.p", "rb"))
fr_embeddings_subset = pickle.load(open("fr_embeddings.p", "rb"))
```

### Look at the data

- en\_embeddings\_subset: the key is an English word, and the value is a 300 dimensional array, which is the embedding for that word.

```
'the': array([ 0.08007812,  0.10498047,  0.04980469,  0.0534668 , -0.06738281, ....
```

- fr\_embeddings\_subset: the key is an French word, and the value is a 300 dimensional array, which is the embedding for that word.

```
'la': array([-6.18250e-03, -9.43867e-04, -8.82648e-03,  3.24623e-02,...
```

### Load two dictionaries mapping the English to French words

- A training dictionary
- and a testing dictionary.

```
In [4]: # Loading the english to french dictionaries
en_fr_train = get_dict('en-fr.train.txt')
print('The length of the English to French training dictionary is', len(en_fr_train))
en_fr_test = get_dict('en-fr.test.txt')
print('The length of the English to French test dictionary is', len(en_fr_train))
```

The length of the English to French training dictionary is 5000  
The length of the English to French test dictionary is 5000

### Looking at the English French dictionary

- en\_fr\_train is a dictionary where the key is the English word and the value is the French translation of that English word.

```
{'the': 'la',
 'and': 'et',
 'was': 'était',
 'for': 'pour',
```

- en\_fr\_test is similar to en\_fr\_train , but is a test set. We won't look at it until we get to testing.

## 1.1 Generate embedding and transform matrices

### Exercise 01: Translating English dictionary to French by using embeddings

You will now implement a function `get_matrices` , which takes the loaded data and returns matrices X and Y .

Inputs:

- `en_fr` : English to French dictionary
- `en_embeddings` : English to embeddings dictionary
- `fr_embeddings` : French to embeddings dictionary

Returns:

- Matrix `X` and matrix `Y`, where each row in `X` is the word embedding for an english word, and the same row in `Y` is the word embedding for the French version of that English word.

X	Y																																																																								
English word embedding { <table border="1" style="border-collapse: collapse; text-align: center;"> <tr><td>2</td><td>9</td><td>...</td><td>...</td><td>1</td><td>3</td></tr> <tr><td>1</td><td>2</td><td>...</td><td>...</td><td>0</td><td>3</td></tr> <tr><td>9</td><td>8</td><td>...</td><td>...</td><td>7</td><td>5</td></tr> <tr><td>1</td><td>0</td><td>...</td><td>...</td><td>9</td><td>3</td></tr> <tr><td>4</td><td>9</td><td>...</td><td>...</td><td>4</td><td>0</td></tr> <tr><td>2</td><td>1</td><td>...</td><td>...</td><td>0</td><td>5</td></tr> </table>	2	9	...	...	1	3	1	2	...	...	0	3	9	8	...	...	7	5	1	0	...	...	9	3	4	9	...	...	4	0	2	1	...	...	0	5	French word embedding { <table border="1" style="border-collapse: collapse; text-align: center;"> <tr><td>1</td><td>9</td><td>...</td><td>...</td><td>9</td><td>3</td></tr> <tr><td>4</td><td>2</td><td>...</td><td>...</td><td>0</td><td>8</td></tr> <tr><td>9</td><td>6</td><td>...</td><td>...</td><td>6</td><td>5</td></tr> <tr><td>1</td><td>9</td><td>...</td><td>...</td><td>9</td><td>3</td></tr> <tr><td>4</td><td>2</td><td>...</td><td>...</td><td>4</td><td>0</td></tr> <tr><td>2</td><td>8</td><td>...</td><td>...</td><td>6</td><td>5</td></tr> </table>	1	9	...	...	9	3	4	2	...	...	0	8	9	6	...	...	6	5	1	9	...	...	9	3	4	2	...	...	4	0	2	8	...	...	6	5
2	9	...	...	1	3																																																																				
1	2	...	...	0	3																																																																				
9	8	...	...	7	5																																																																				
1	0	...	...	9	3																																																																				
4	9	...	...	4	0																																																																				
2	1	...	...	0	5																																																																				
1	9	...	...	9	3																																																																				
4	2	...	...	0	8																																																																				
9	6	...	...	6	5																																																																				
1	9	...	...	9	3																																																																				
4	2	...	...	4	0																																																																				
2	8	...	...	6	5																																																																				

Figure 2

Use the `en_fr` dictionary to ensure that the ith row in the `X` matrix corresponds to the ith row in the `Y` matrix.

**Instructions:** Complete the function `get_matrices()`:

- Iterate over English words in `en_fr` dictionary.
- Check if the word have both English and French embedding.

### ► Hints

```
In [5]: # UNQ_C1 (UNIQUE CELL IDENTIFIER, DO NOT EDIT)
def get_matrices(en_fr, french_vecs, english_vecs):
    """
    Input:
        en_fr: English to French dictionary
        french_vecs: French words to their corresponding word embeddings.
        english_vecs: English words to their corresponding word embeddings.
    Output:
        X: a matrix where the columns are the English embeddings.
        Y: a matrix where the columns correspond to the French embeddings.
        R: the projection matrix that minimizes the F norm ||X R -Y||^2.
    """

    ### START CODE HERE (REPLACE INSTANCES OF 'None' with your code) ###

    # X_L and Y_L are lists of the english and french word embeddings
    X_1 = list()
    Y_1 = list()

    # get the english words (the keys in the dictionary) and store in a set()
    english_set = english_vecs.keys()

    # get the french words (keys in the dictionary) and store in a set()
    french_set = french_vecs.keys()

    # store the french words that are part of the english-french dictionary (these)
    french_words = set(en_fr.values())

    # Loop through all english, french word pairs in the english-french dictionary
    for en_word, fr_word in en_fr.items():

        # check that the french word has an embedding and that the english word is
        if fr_word in french_set and en_word in english_set:

            # get the english embedding
            en_vec = english_vecs[en_word]

            # get the french embedding
            fr_vec = french_vecs[fr_word]

            # add the english embedding to the list
            X_1.append(en_vec)

            # add the french embedding to the list
            Y_1.append(fr_vec)

    # stack the vectors of X_L into a matrix X
    X = np.vstack(X_1)

    # stack the vectors of Y_L into a matrix Y
    Y = np.vstack(Y_1)

    ### END CODE HERE ###


    return X, Y
```

Now we will use function `get_matrices()` to obtain sets `X_train` and `Y_train` of English and French word embeddings into the corresponding vector space models.


```
In [6]: # UNQ_C2 (UNIQUE CELL IDENTIFIER, DO NOT EDIT)
# You do not have to input any code in this cell, but it is relevant to grading,
# getting the training set:
X_train, Y_train = get_matrices(
    en_fr_train, fr_embeddings_subset, en_embeddings_subset)
```

## 2. Translations

English



Francais



Natural Language Processing



Figure 1

Write a program that translates English words to French words using word embeddings and vector space models.

### 2.1 Translation as linear transformation of embeddings

Given dictionaries of English and French word embeddings you will create a transformation matrix  $R$

- Given an English word embedding,  $e$ , you can multiply  $eR$  to get a new word embedding  $f$ .
  - Both  $e$  and  $f$  are [row vectors](https://en.wikipedia.org/wiki/Row_and_column_vectors) ([https://en.wikipedia.org/wiki/Row\\_and\\_column\\_vectors](https://en.wikipedia.org/wiki/Row_and_column_vectors)).
- You can then compute the nearest neighbors to  $f$  in the french embeddings and recommend the word that is most similar to the transformed word embedding.

### Describing translation as the minimization problem

Find a matrix  $R$  that minimizes the following equation.

$$\arg \min_R \|XR - Y\|_F \quad (1)$$

### Frobenius norm

The Frobenius norm of a matrix  $A$  (assuming it is of dimension  $m, n$ ) is defined as the square root of the sum of the absolute squares of its elements:

$$\|A\|_F \equiv \sqrt{\sum_{i=1}^m \sum_{j=1}^n |a_{ij}|^2} \quad (2)$$

## Actual loss function

In the real world applications, the Frobenius norm loss:

$$\|\mathbf{XR} - \mathbf{Y}\|_F$$

is often replaced by its squared value divided by  $m$ :

$$\frac{1}{m} \|\mathbf{XR} - \mathbf{Y}\|_F^2$$

where  $m$  is the number of examples (rows in  $\mathbf{X}$ ).

- The same  $R$  is found when using this loss function versus the original Frobenius norm.
- The reason for taking the square is that it's easier to compute the gradient of the squared Frobenius.
- The reason for dividing by  $m$  is that we're more interested in the average loss per embedding than the loss for the entire training set.
  - The loss for all training set increases with more words (training examples), so taking the average helps us to track the average loss regardless of the size of the training set.

**[Optional] Detailed explanation why we use norm squared instead of the norm:**

- ▶ Click for optional details

## Exercise 02: Implementing translation mechanism described in this section.

### Step 1: Computing the loss

- The loss function will be squared Frobenius norm of the difference between matrix and its approximation, divided by the number of training examples  $m$ .
- Its formula is:

$$L(X, Y, R) = \frac{1}{m} \sum_{i=1}^m \sum_{j=1}^n (a_{ij})^2$$

where  $a_{ij}$  is value in  $i$ th row and  $j$ th column of the matrix  $\mathbf{XR} - \mathbf{Y}$ .

**Instructions: complete the `compute_loss()` function**

- Compute the approximation of  $Y$  by matrix multiplying  $X$  and  $R$

- Compute difference  $XR - Y$
- Compute the squared Frobenius norm of the difference and divide it by  $m$ .

## ► Hints

```
In [7]: # UNQ_C3 (UNIQUE CELL IDENTIFIER, DO NOT EDIT)
def compute_loss(X, Y, R):
    """
    Inputs:
        X: a matrix of dimension (m,n) where the columns are the English embeddings
        Y: a matrix of dimension (m,n) where the columns correspond to the French words
        R: a matrix of dimension (n,n) - transformation matrix from English to French words
    Outputs:
        L: a matrix of dimension (m,n) - the value of the loss function for given X, Y, R
    ...
    ### START CODE HERE (REPLACE INSTANCES OF 'None' with your code) #####
    # m is the number of rows in X
    m = len(X)

    # diff is XR - Y
    diff = np.dot(X,R) - Y

    # diff_squared is the element-wise square of the difference
    diff_squared = np.square(diff)

    # sum_diff_squared is the sum of the squared elements
    sum_diff_squared = np.sum(diff_squared)

    # loss is the sum_diff_squared divided by the number of examples (m)
    loss = sum_diff_squared/m
    ### END CODE HERE #####
    return loss
```

## Exercise 03

### Step 2: Computing the gradient of loss in respect to transform matrix R

- Calculate the gradient of the loss with respect to transform matrix  $R$ .
- The gradient is a matrix that encodes how much a small change in  $R$  affect the change in the loss function.
- The gradient gives us the direction in which we should decrease  $R$  to minimize the loss.
- $m$  is the number of training examples (number of rows in  $X$ ).
- The formula for the gradient of the loss function  $L(X, Y, R)$  is:

$$\frac{d}{dR} L(X, Y, R) = \frac{d}{dR} \left( \frac{1}{m} \|XR - Y\|_F^2 \right) = \frac{2}{m} X^T(XR - Y)$$

**Instructions:** Complete the `compute_gradient` function below.

## ► Hints

```
In [12]: # UNQ_C4 (UNIQUE CELL IDENTIFIER, DO NOT EDIT)
def compute_gradient(X, Y, R):
    ...
    Inputs:
        X: a matrix of dimension (m,n) where the columns are the English embedding
        Y: a matrix of dimension (m,n) where the columns correspond to the French
        R: a matrix of dimension (n,n) - transformation matrix from English to French
    Outputs:
        g: a matrix of dimension (n,n) - gradient of the loss function L for given R
    ...
    ### START CODE HERE (REPLACE INSTANCES OF 'None' with your code) #####
    # m is the number of rows in X
    m = len(X)

    # gradient is  $X^T(XR - Y) * 2/m$ 
    gradient = np.dot(np.transpose(X), np.dot(X, R) - Y) * 2/m
    ##### END CODE HERE #####
    return gradient
```

## Step 3: Finding the optimal $R$ with gradient descent algorithm

### Gradient descent

Gradient descent ([https://ml-cheatsheet.readthedocs.io/en/latest/gradient\\_descent.html](https://ml-cheatsheet.readthedocs.io/en/latest/gradient_descent.html)) is an iterative algorithm which is used in searching for the optimum of the function.

- Earlier, we've mentioned that the gradient of the loss with respect to the matrix encodes how much a tiny change in some coordinate of that matrix affect the change of loss function.
- Gradient descent uses that information to iteratively change matrix  $R$  until we reach a point where the loss is minimized.

### Training with a fixed number of iterations

Most of the time we iterate for a fixed number of training steps rather than iterating until the loss falls below a threshold.

**OPTIONAL: explanation for fixed number of iterations**

► [click here for detailed discussion](#)

Pseudocode:

1. Calculate gradient  $g$  of the loss with respect to the matrix  $R$ .
2. Update  $R$  with the formula:

$$R_{\text{new}} = R_{\text{old}} - \alpha g$$

Where  $\alpha$  is the learning rate, which is a scalar.

## Learning rate

- The learning rate or "step size"  $\alpha$  is a coefficient which decides how much we want to change  $R$  in each step.
- If we change  $R$  too much, we could skip the optimum by taking too large of a step.
- If we make only small changes to  $R$ , we will need many steps to reach the optimum.
- Learning rate  $\alpha$  is used to control those changes.
- Values of  $\alpha$  are chosen depending on the problem, and we'll use `learning_rate = 0.0003` as the default value for our algorithm.

## Exercise 04

**Instructions:** Implement `align_embeddings()`

### ► Hints

```
In [15]: # UNQ_C5 (UNIQUE CELL IDENTIFIER, DO NOT EDIT)
def align_embeddings(X, Y, train_steps=100, learning_rate=0.0003):
    """
    Inputs:
        X: a matrix of dimension (m,n) where the columns are the English embeddings
        Y: a matrix of dimension (m,n) where the columns correspond to the French embeddings
        train_steps: positive int - describes how many steps will gradient descent take
        learning_rate: positive float - describes how big steps will gradient descent take
    Outputs:
        R: a matrix of dimension (n,n) - the projection matrix that minimizes the loss
    ...
    np.random.seed(129)

    # the number of columns in X is the number of dimensions for a word vector (e.g., 300)
    # R is a square matrix with length equal to the number of dimensions in the vectors (e.g., 300)
    R = np.random.rand(X.shape[1], X.shape[1])

    for i in range(train_steps):
        if i % 25 == 0:
            print(f"loss at iteration {i} is: {compute_loss(X, Y, R):.4f}")
            ### START CODE HERE (REPLACE INSTANCES OF 'None' with your code) #####
            # use the function that you defined to compute the gradient
            gradient = compute_gradient(X, Y, R)

            # update R by subtracting the Learning rate times gradient
            R -= learning_rate * gradient
            ### END CODE HERE #####
    return R
```

```
In [16]: # UNQ_C6 (UNIQUE CELL IDENTIFIER, DO NOT EDIT)
# You do not have to input any code in this cell, but it is relevant to grading,
# Testing your implementation.
np.random.seed(129)
m = 10
n = 5
X = np.random.rand(m, n)
Y = np.random.rand(m, n) * .1
R = align_embeddings(X, Y)
```

```
loss at iteration 0 is: 3.7242
loss at iteration 25 is: 3.6283
loss at iteration 50 is: 3.5350
loss at iteration 75 is: 3.4442
```

### Expected Output:

```
loss at iteration 0 is: 3.7242
loss at iteration 25 is: 3.6283
loss at iteration 50 is: 3.5350
loss at iteration 75 is: 3.4442
```

## Calculate transformation matrix R

Using those the training set, find the transformation matrix **R** by calling the function `align_embeddings()`.

**NOTE:** The code cell below will take a few minutes to fully execute (~3 mins)

```
In [17]: # UNQ_C7 (UNIQUE CELL IDENTIFIER, DO NOT EDIT)
# You do not have to input any code in this cell, but it is relevant to grading,
R_train = align_embeddings(X_train, Y_train, train_steps=400, learning_rate=0.8)
```

```
loss at iteration 0 is: 963.0146
loss at iteration 25 is: 97.8292
loss at iteration 50 is: 26.8329
loss at iteration 75 is: 9.7893
loss at iteration 100 is: 4.3776
loss at iteration 125 is: 2.3281
loss at iteration 150 is: 1.4480
loss at iteration 175 is: 1.0338
loss at iteration 200 is: 0.8251
loss at iteration 225 is: 0.7145
loss at iteration 250 is: 0.6534
loss at iteration 275 is: 0.6185
loss at iteration 300 is: 0.5981
loss at iteration 325 is: 0.5858
loss at iteration 350 is: 0.5782
loss at iteration 375 is: 0.5735
```

### Expected Output

```

loss at iteration 0 is: 963.0146
loss at iteration 25 is: 97.8292
loss at iteration 50 is: 26.8329
loss at iteration 75 is: 9.7893
loss at iteration 100 is: 4.3776
loss at iteration 125 is: 2.3281
loss at iteration 150 is: 1.4480
loss at iteration 175 is: 1.0338
loss at iteration 200 is: 0.8251
loss at iteration 225 is: 0.7145
loss at iteration 250 is: 0.6534
loss at iteration 275 is: 0.6185
loss at iteration 300 is: 0.5981
loss at iteration 325 is: 0.5858
loss at iteration 350 is: 0.5782
loss at iteration 375 is: 0.5735

```

## 2.2 Testing the translation

### k-Nearest neighbors algorithm

[k-Nearest neighbors algorithm \(\[https://en.wikipedia.org/wiki/K-nearest\\\_neighbors\\\_algorithm\]\(https://en.wikipedia.org/wiki/K-nearest\_neighbors\_algorithm\)\)](https://en.wikipedia.org/wiki/K-nearest_neighbors_algorithm)

- k-NN is a method which takes a vector as input and finds the other vectors in the dataset that are closest to it.
- The 'k' is the number of "nearest neighbors" to find (e.g. k=2 finds the closest two neighbors).

### Searching for the translation embedding

Since we're approximating the translation function from English to French embeddings by a linear transformation matrix  $\mathbf{R}$ , most of the time we won't get the exact embedding of a French word when we transform embedding  $\mathbf{e}$  of some particular English word into the French embedding space.

- This is where  $k$ -NN becomes really useful! By using 1-NN with  $\mathbf{eR}$  as input, we can search for an embedding  $\mathbf{f}$  (as a row) in the matrix  $\mathbf{Y}$  which is the closest to the transformed vector  $\mathbf{eR}$

### Cosine similarity

Cosine similarity between vectors  $u$  and  $v$  calculated as the cosine of the angle between them.  
The formula is

$$\cos(u, v) = \frac{u \cdot v}{\|u\| \|v\|}$$

- $\cos(u, v) = 1$  when  $u$  and  $v$  lie on the same line and have the same direction.
- $\cos(u, v)$  is  $-1$  when they have exactly opposite directions.
- $\cos(u, v)$  is  $0$  when the vectors are orthogonal (perpendicular) to each other.

**Note:** Distance and similarity are pretty much opposite things.

- We can obtain distance metric from cosine similarity, but the cosine similarity can't be used directly as the distance metric.
- When the cosine similarity increases (towards 1), the "distance" between the two vectors decreases (towards 0).
- We can define the cosine distance between  $u$  and  $v$  as

$$d_{\cos}(u, v) = 1 - \cos(u, v)$$

**Exercise 05:** Complete the function `nearest_neighbor()`

Inputs:

- Vector  $v$ ,
- A set of possible nearest neighbors `candidates`
- $k$  nearest neighbors to find.
- The distance metric should be based on cosine similarity.
- `cosine_similarity` function is already implemented and imported for you. It's arguments are two vectors and it returns the cosine of the angle between them.
- Iterate over rows in `candidates`, and save the result of similarities between current row and vector  $v$  in a python list. Take care that similarities are in the same order as row vectors of `candidates`.
- Now you can use `numpy.argsort` (<https://docs.scipy.org/doc/numpy/reference/generated/numpy.argsort.html#numpy.argsort>) to sort the indices for the rows of `candidates`.

## ► Hints

```
In [20]: # UNQ_C8 (UNIQUE CELL IDENTIFIER, DO NOT EDIT)
def nearest_neighbor(v, candidates, k=1):
    """
    Input:
        - v, the vector you are going to find the nearest neighbor for
        - candidates: a set of vectors where we will find the neighbors
        - k: top k nearest neighbors to find
    Output:
        - k_idx: the indices of the top k closest vectors in sorted form
    """
    ### START CODE HERE (REPLACE INSTANCES OF 'None' with your code) ####
    similarity_l = []

    # for each candidate vector...
    for row in candidates:
        # get the cosine similarity
        cos_similarity = cosine_similarity(v, row)

        # append the similarity to the list
        similarity_l.append(cos_similarity)

    # sort the similarity list and get the indices of the sorted list
    sorted_ids = np.argsort(similarity_l)

    # get the indices of the k most similar candidate vectors
    k_idx = sorted_ids[-k:]
    ### END CODE HERE ####
    return k_idx
```

```
In [21]: # UNQ_C9 (UNIQUE CELL IDENTIFIER, DO NOT EDIT)
# You do not have to input any code in this cell, but it is relevant to grading,
# Test your implementation:
v = np.array([1, 0, 1])
candidates = np.array([[1, 0, 5], [-2, 5, 3], [2, 0, 1], [6, -9, 5], [9, 9, 9]])
print(candidates[nearest_neighbor(v, candidates, 3)])
```

```
[[9 9 9]
 [1 0 5]
 [2 0 1]]
```

**Expected Output:**

```
[[9 9 9]
 [1 0 5]
 [2 0 1]]
```

## Test your translation and compute its accuracy

**Exercise 06:** Complete the function `test_vocabulary` which takes in English embedding matrix  $X$ , French embedding matrix  $Y$  and the  $R$  matrix and returns the accuracy of translations from  $X$  to  $Y$  by  $R$ .

- Iterate over transformed English word embeddings and check if the closest French word vector belongs to French word that is the actual translation.
- Obtain an index of the closest French embedding by using `nearest_neighbor` (with argument `k=1`), and compare it to the index of the English embedding you have just transformed.
- Keep track of the number of times you get the correct translation.
- Calculate accuracy as

$$\text{accuracy} = \frac{\#\text{(correct predictions)}}{\#\text{(total predictions)}}$$

```
In [22]: # UNQ_C10 (UNIQUE CELL IDENTIFIER, DO NOT EDIT)
def test_vocabulary(X, Y, R):
    ...

    Input:
        X: a matrix where the columns are the English embeddings.
        Y: a matrix where the columns correspond to the French embeddings.
        R: the transform matrix which translates word embeddings from
            English to French word vector space.

    Output:
        accuracy: for the English to French capitals
    ...

    ### START CODE HERE (REPLACE INSTANCES OF 'None' with your code) ###
    # The prediction is X times R
    pred = np.dot(X,R)

    # initialize the number correct to zero
    num_correct = 0

    # Loop through each row in pred (each transformed embedding)
    for i in range(len(pred)):
        # get the index of the nearest neighbor of pred at row 'i'; also pass in
        pred_idx = nearest_neighbor(pred[i], Y)

        # if the index of the nearest neighbor equals the row of i...
        if pred_idx == i:
            # increment the number correct by 1.
            num_correct += 1

    # accuracy is the number correct divided by the number of rows in 'pred' (also
    accuracy = num_correct/len(X)

    ### END CODE HERE ###

    return accuracy
```

Let's see how is your translation mechanism working on the unseen data:

```
In [23]: X_val, Y_val = get_matrices(en_fr_test, fr_embeddings_subset, en_embeddings_subset)
```

```
In [24]: # UNQ_C11 (UNIQUE CELL IDENTIFIER, DO NOT EDIT)
# You do not have to input any code in this cell, but it is relevant to grading,
acc = test_vocabulary(X_val, Y_val, R_train) # this might take a minute or two
print(f"accuracy on test set is {acc:.3f}")
```

accuracy on test set is 0.557

#### Expected Output:

0.557

You managed to translate words from one language to another language without ever seeing them with almost 56% accuracy by using some basic linear algebra and learning a mapping of words from one language to another!

## 3. LSH and document search

In this part of the assignment, you will implement a more efficient version of k-nearest neighbors using locality sensitive hashing. You will then apply this to document search.

- Process the tweets and represent each tweet as a vector (represent a document with a vector embedding).
- Use locality sensitive hashing and k nearest neighbors to find tweets that are similar to a given tweet.

```
In [25]: # get the positive and negative tweets
all_positive_tweets = twitter_samples.strings('positive_tweets.json')
all_negative_tweets = twitter_samples.strings('negative_tweets.json')
all_tweets = all_positive_tweets + all_negative_tweets
```

### 3.1 Getting the document embeddings

#### Bag-of-words (BOW) document models

Text documents are sequences of words.

- The ordering of words makes a difference. For example, sentences "Apple pie is better than pepperoni pizza." and "Pepperoni pizza is better than apple pie" have opposite meanings due to the word ordering.
- However, for some applications, ignoring the order of words can allow us to train an efficient and still effective model.
- This approach is called Bag-of-words document model.

#### Document embeddings

- Document embedding is created by summing up the embeddings of all words in the document.
- If we don't know the embedding of some word, we can ignore that word.

**Exercise 07:** Complete the `get_document_embedding()` function.

- The function `get_document_embedding()` encodes entire document as a "document" embedding.
- It takes in a document (as a string) and a dictionary, `en_embeddings`
- It processes the document, and looks up the corresponding embedding of each word.
- It then sums them up and returns the sum of all word vectors of that processed tweet.

## ► Hints

```
In [26]: # UNQ_C12 (UNIQUE CELL IDENTIFIER, DO NOT EDIT)
def get_document_embedding(tweet, en_embeddings):
    """
    Input:
        - tweet: a string
        - en_embeddings: a dictionary of word embeddings
    Output:
        - doc_embedding: sum of all word embeddings in the tweet
    ...
    doc_embedding = np.zeros(300)

    ### START CODE HERE (REPLACE INSTANCES OF 'None' with your code) ####
    # process the document into a list of words (process the tweet)
    processed_doc = process_tweet(tweet)
    for word in processed_doc:
        # add the word embedding to the running total for the document embedding
        doc_embedding = doc_embedding + en_embeddings.get(word, [0])
    ### END CODE HERE ####
    return doc_embedding
```

```
In [27]: # UNQ_C13 (UNIQUE CELL IDENTIFIER, DO NOT EDIT)
# You do not have to input any code in this cell, but it is relevant to grading,
# testing your function
custom_tweet = "RT @Twitter @chapagain Hello There! Have a great day. :)" #good tweet
tweet_embedding = get_document_embedding(custom_tweet, en_embeddings_subset)
tweet_embedding[-5:]
```

Out[27]: array([-0.00268555, -0.15378189, -0.55761719, -0.07216644, -0.32263184])

**Expected output:**

array([-0.00268555, -0.15378189, -0.55761719, -0.07216644, -0.32263184])

## Exercise 08

### Store all document vectors into a dictionary

Now, let's store all the tweet embeddings into a dictionary. Implement `get_document_vecs()`

```
In [28]: # UNQ_C14 (UNIQUE CELL IDENTIFIER, DO NOT EDIT)
def get_document_vecs(all_docs, en_embeddings):
    ...

    Input:
        - all_docs: list of strings - all tweets in our dataset.
        - en_embeddings: dictionary with words as the keys and their embeddings as values.

    Output:
        - document_vec_matrix: matrix of tweet embeddings.
        - ind2Doc_dict: dictionary with indices of tweets in vecs as keys and their document embeddings as values.

    ...

    # the dictionary's key is an index (integer) that identifies a specific tweet
    # the value is the document embedding for that document
    ind2Doc_dict = {}

    # this is list that will store the document vectors
    document_vec_1 = []

    for i, doc in enumerate(all_docs):

        ### START CODE HERE (REPLACE INSTANCES OF 'None' with your code) #####
        # get the document embedding of the tweet
        doc_embedding = get_document_embedding(doc, en_embeddings)

        # save the document embedding into the ind2Tweet dictionary at index i
        ind2Doc_dict[i] = doc_embedding

        # append the document embedding to the list of document vectors
        document_vec_1.append(doc_embedding)

        ### END CODE HERE #####
        # convert the list of document vectors into a 2D array (each row is a document)
        document_vec_matrix = np.vstack(document_vec_1)

    return document_vec_matrix, ind2Doc_dict
```

```
In [29]: document_vecs, ind2Tweet = get_document_vecs(all_tweets, en_embeddings_subset)
```

```
In [30]: # UNQ_C15 (UNIQUE CELL IDENTIFIER, DO NOT EDIT)
# You do not have to input any code in this cell, but it is relevant to grading.

print(f"length of dictionary {len(ind2Tweet)}")
print(f"shape of document_vecs {document_vecs.shape}")

length of dictionary 10000
shape of document_vecs (10000, 300)
```

***Expected Output***

```
length of dictionary 10000
shape of document_vecs (10000, 300)
```

## 3.2 Looking up the tweets

Now you have a vector of dimension (m,d) where `m` is the number of tweets (10,000) and `d` is the dimension of the embeddings (300). Now you will input a tweet, and use cosine similarity to see which tweet in our corpus is similar to your tweet.

```
In [31]: my_tweet = 'i am sad'
process_tweet(my_tweet)
tweet_embedding = get_document_embedding(my_tweet, en_embeddings_subset)
```

```
In [32]: # UNQ_C16 (UNIQUE CELL IDENTIFIER, DO NOT EDIT)
# You do not have to input any code in this cell, but it is relevant to grading,
# this gives you a similar tweet as your input.
# this implementation is vectorized...
idx = np.argmax(cosine_similarity(document_vecs, tweet_embedding))
print(all_tweets[idx])
```

```
@zoeeylim sad sad sad kid :( it's ok I help you watch the match HAHAHAHAHA
```

***Expected Output***

```
@zoeeylim sad sad sad kid :( it's ok I help you watch the match HAHAHAHA
HA
```

## 3.3 Finding the most similar tweets with LSH

You will now implement locality sensitive hashing (LSH) to identify the most similar tweet.

- Instead of looking at all 10,000 vectors, you can just search a subset to find its nearest neighbors.

Let's say your data points are plotted like this:




Figure 3

You can divide the vector space into regions and search within one region for nearest neighbors of a given vector.




Figure 4

```
In [33]: N_VECS = len(all_tweets)      # This many vectors.
N_DIMS = len(ind2Tweet[1])        # Vector dimensionality.
print(f"Number of vectors is {N_VECS} and each has {N_DIMS} dimensions.")
```

Number of vectors is 10000 and each has 300 dimensions.

### Choosing the number of planes

- Each plane divides the space to 2 parts.
- So  $n$  planes divide the space into  $2^n$  hash buckets.
- We want to organize 10,000 document vectors into buckets so that every bucket has about 16 vectors.
- For that we need  $\frac{10000}{16} = 625$  buckets.
- We're interested in  $n$ , number of planes, so that  $2^n = 625$ . Now, we can calculate  $n = \log_2 625 = 9.29 \approx 10$ .

```
In [34]: # The number of planes. We use Log2(625) to have ~16 vectors/bucket.
N_PLANES = 10
# Number of times to repeat the hashing to improve the search.
N_UNIVERSES = 25
```

## 3.4 Getting the hash number for a vector

For each vector, we need to get a unique number associated to that vector in order to assign it to a "hash bucket".

### Hyperplanes in vector spaces

- In 3-dimensional vector space, the hyperplane is a regular plane. In 2 dimensional vector space, the hyperplane is a line.
- Generally, the hyperplane is subspace which has dimension 1 lower than the original vector space has.
- A hyperplane is uniquely defined by its normal vector.
- Normal vector  $n$  of the plane  $\pi$  is the vector to which all vectors in the plane  $\pi$  are orthogonal (perpendicular in 3 dimensional case).

### Using Hyperplanes to split the vector space

We can use a hyperplane to split the vector space into 2 parts.

- All vectors whose dot product with a plane's normal vector is positive are on one side of the plane.
- All vectors whose dot product with the plane's normal vector is negative are on the other side of the plane.

### Encoding hash buckets

- For a vector, we can take its dot product with all the planes, then encode this information to assign the vector to a single hash bucket.
- When the vector is pointing to the opposite side of the hyperplane than normal, encode it by 0.
- Otherwise, if the vector is on the same side as the normal vector, encode it by 1.
- If you calculate the dot product with each plane in the same order for every vector, you've encoded each vector's unique hash ID as a binary number, like [0, 1, 1, ... 0].

### Exercise 09: Implementing hash buckets

We've initialized hash table `hashes` for you. It is list of `N_UNIVERSES` matrices, each describes its own hash table. Each matrix has `N_DIMS` rows and `N_PLANES` columns. Every column of that matrix is a `N_DIMS`-dimensional normal vector for each of `N_PLANES` hyperplanes which are used for creating buckets of the particular hash table.

*Exercise:* Your task is to complete the function `hash_value_of_vector` which places vector `v` in the correct hash bucket.

- First multiply your vector `v`, with a corresponding plane. This will give you a vector of dimension  $(1, N\_planes)$ .
- You will then convert every element in that vector to 0 or 1.
- You create a hash vector by doing the following: if the element is negative, it becomes a 0, otherwise you change it to a 1.

- You then compute the unique number for the vector by iterating over `N_PLANES`
- Then you multiply  $2^i$  times the corresponding bit (0 or 1).
- You will then store that sum in the variable `hash_value`.

**Instructions:** Create a hash for the vector in the function below. Use this formula:

$$\text{hash} = \sum_{i=0}^{N-1} (2^i \times h_i)$$

### Create the sets of planes

- Create multiple (25) sets of planes (the planes that divide up the region).
- You can think of these as 25 separate ways of dividing up the vector space with a different set of planes.
- Each element of this list contains a matrix with 300 rows (the word vector have 300 dimensions), and 10 columns (there are 10 planes in each "universe").

```
In [35]: np.random.seed(0)
planes_1 = [np.random.normal(size=(N_DIMS, N_PLANES))
            for _ in range(N_UNIVERSES)]
```

### ► Hints

```
In [36]: # UNQ_C17 (UNIQUE CELL IDENTIFIER, DO NOT EDIT)
def hash_value_of_vector(v, planes):
    """Create a hash for a vector; hash_id says which random hash to use.
    Input:
        - v: vector of tweet. It's dimension is (1, N_DIMS)
        - planes: matrix of dimension (N_DIMS, N_PLANES) - the set of planes that
    Output:
        - res: a number which is used as a hash for your vector
    """
    ### START CODE HERE (REPLACE INSTANCES OF 'None' with your code) ###
    # for the set of planes,
    # calculate the dot product between the vector and the matrix containing the
    # remember that planes has shape (300, 10)
    # The dot product will have the shape (1,10)
    dot_product = np.dot(v, planes)

    # get the sign of the dot product (1,10) shaped vector
    sign_of_dot_product = np.sign(dot_product)

    # set h to be false (equivalent to 0 when used in operations) if the sign is -
    # and true (equivalent to 1) if the sign is positive (1,10) shaped vector
    h = sign_of_dot_product >= 0

    # remove extra un-used dimensions (convert this from a 2D to a 1D array)
    h = np.squeeze(h)

    # initialize the hash value to 0
    hash_value = 0

    n_planes = planes.shape[1]
    for i in range(n_planes):
        # increment the hash value by 2^i * h_i
        hash_value += (2**i) * h[i]
    ### END CODE HERE ###

    # cast hash_value as an integer
    hash_value = int(hash_value)

    return hash_value
```

```
In [37]: # UNQ_C18 (UNIQUE CELL IDENTIFIER, DO NOT EDIT)
# You do not have to input any code in this cell, but it is relevant to grading,

np.random.seed(0)
idx = 0
planes = planes_l[idx] # get one 'universe' of planes to test the function
vec = np.random.rand(1, 300)
print(f" The hash value for this vector,",
      f"and the set of planes at index {idx},",
      f"is {hash_value_of_vector(vec, planes)}")
```

The hash value for this vector, and the set of planes at index 0, is 768


**Expected Output**

The hash value for this vector, and the set of planes at index 0, is 768

## 3.5 Creating a hash table

### Exercise 10

Given that you have a unique number for each vector (or tweet), You now want to create a hash table. You need a hash table, so that given a hash\_id, you can quickly look up the corresponding vectors. This allows you to reduce your search by a significant amount of time.



We have given you the `make_hash_table` function, which maps the tweet vectors to a bucket and stores the vector there. It returns the `hash_table` and the `id_table`. The `id_table` allows you know which vector in a certain bucket corresponds to what tweet.

► **Hints**

```
In [38]: # UNQ_C19 (UNIQUE CELL IDENTIFIER, DO NOT EDIT)
# This is the code used to create a hash table: feel free to read over it
def make_hash_table(vecs, planes):
    """
    Input:
        - vecs: list of vectors to be hashed.
        - planes: the matrix of planes in a single "universe", with shape (embedd...
    Output:
        - hash_table: dictionary - keys are hashes, values are lists of vectors
        - id_table: dictionary - keys are hashes, values are list of vectors id's
                      (it's used to know which tweet corresponds to the hash)
    """
    ### START CODE HERE (REPLACE INSTANCES OF 'None' with your code) ###

    # number of planes is the number of columns in the planes matrix
    num_of_planes = len(planes[0])

    # number of buckets is 2^(number of planes)
    num_buckets = 2 ** num_of_planes

    # create the hash table as a dictionary.
    # Keys are integers (0,1,2... number of buckets)
    # Values are empty lists
    hash_table = { i : [] for i in range(num_buckets) }

    # create the id table as a dictionary.
    # Keys are integers (0,1,2... number of buckets)
    # Values are empty lists
    id_table = { i : [] for i in range(num_buckets) }

    # for each vector in 'vecs'
    for i, v in enumerate(vecs):
        # calculate the hash value for the vector
        h = hash_value_of_vector(v, planes)

        # store the vector into hash_table at key h,
        # by appending the vector v to the list at key h
        hash_table[h].append(v)

        # store the vector's index 'i' (each document is given a unique integer)
        # the key is the h, and the 'i' is appended to the list at key h
        id_table[h].append(i)

    ### END CODE HERE ###

    return hash_table, id_table
```

```
In [39]: # UNQ_C20 (UNIQUE CELL IDENTIFIER, DO NOT EDIT)
# You do not have to input any code in this cell, but it is relevant to grading,
np.random.seed(0)
planes = planes_1[0] # get one 'universe' of planes to test the function
vec = np.random.rand(1, 300)
tmp_hash_table, tmp_id_table = make_hash_table(document_vecs, planes)

print(f"The hash table at key 0 has {len(tmp_hash_table[0])} document vectors")
print(f"The id table at key 0 has {len(tmp_id_table[0])}")
print(f"The first 5 document indices stored at key 0 of are {tmp_id_table[0][0:5])
```

```
The hash table at key 0 has 3 document vectors
The id table at key 0 has 3
The first 5 document indices stored at key 0 of are [3276, 3281, 3282]
```

### Expected output

```
The hash table at key 0 has 3 document vectors
The id table at key 0 has 3
The first 5 document indices stored at key 0 of are [3276, 3281, 3282]
```

## 3.6 Creating all hash tables

You can now hash your vectors and store them in a hash table that would allow you to quickly look up and search for similar vectors. Run the cell below to create the hashes. By doing so, you end up having several tables which have all the vectors. Given a vector, you then identify the buckets in all the tables. You can then iterate over the buckets and consider much fewer vectors. The more buckets you use, the more accurate your lookup will be, but also the longer it will take.

```
In [40]: # Creating the hashtables
hash_tables = []
id_tables = []
for universe_id in range(N_UNIVERSES): # there are 25 hashes
    print('working on hash universe #:', universe_id)
    planes = planes_1[universe_id]
    hash_table, id_table = make_hash_table(document_vecs, planes)
    hash_tables.append(hash_table)
    id_tables.append(id_table)
```

```
working on hash universe #: 0
working on hash universe #: 1
working on hash universe #: 2
working on hash universe #: 3
working on hash universe #: 4
working on hash universe #: 5
working on hash universe #: 6
working on hash universe #: 7
working on hash universe #: 8
working on hash universe #: 9
working on hash universe #: 10
working on hash universe #: 11
working on hash universe #: 12
working on hash universe #: 13
working on hash universe #: 14
working on hash universe #: 15
working on hash universe #: 16
working on hash universe #: 17
working on hash universe #: 18
working on hash universe #: 19
working on hash universe #: 20
working on hash universe #: 21
working on hash universe #: 22
working on hash universe #: 23
working on hash universe #: 24
```

## Approximate K-NN

### Exercise 11

Implement approximate K nearest neighbors using locality sensitive hashing, to search for documents that are similar to a given document at the index `doc_id`.

#### Inputs

- `doc_id` is the index into the document list `all_tweets`.
- `v` is the document vector for the tweet in `all_tweets` at index `doc_id`.
- `planes_1` is the list of planes (the global variable created earlier).
- `k` is the number of nearest neighbors to search for.
- `num_universes_to_use` : to save time, we can use fewer than the total number of available universes. By default, it's set to `N_UNIVERSES`, which is 25 for this assignment.

The `approximate_knn` function finds a subset of candidate vectors that are in the same "hash bucket" as the input vector 'v'. Then it performs the usual k-nearest neighbors search on this subset (instead of searching through all 10,000 tweets).

► **Hints**

```
In [41]: # UNQ_C21 (UNIQUE CELL IDENTIFIER, DO NOT EDIT)
# This is the code used to do the fast nearest neighbor search. Feel free to go over it.
def approximate_knn(doc_id, v, planes_l, k=1, num_universes_to_use=N_UNIVERSES):
    """Search for k-NN using hashes."""
    assert num_universes_to_use <= N_UNIVERSES

    # Vectors that will be checked as possible nearest neighbor
    vecs_to_consider_l = list()

    # List of document IDs
    ids_to_consider_l = list()

    # create a set for ids to consider, for faster checking if a document ID already
    # been considered
    ids_to_consider_set = set()

    # Loop through the universes of planes
    for universe_id in range(num_universes_to_use):

        # get the set of planes from the planes_l list, for this particular universe
        planes = planes_l[universe_id]

        # get the hash value of the vector for this set of planes
        hash_value = hash_value_of_vector(v, planes)

        # get the hash table for this particular universe_id
        hash_table = hash_tables[universe_id]

        # get the list of document vectors for this hash table, where the key is
        # the hash value
        document_vectors_l = hash_table[hash_value]

        # get the id_table for this particular universe_id
        id_table = id_tables[universe_id]

        # get the subset of documents to consider as nearest neighbors from this
        # hash table
        new_ids_to_consider = id_table[hash_value]

        #### START CODE HERE (REPLACE INSTANCES OF 'None' with your code) ####

        # remove the id of the document that we're searching
        if doc_id in new_ids_to_consider:
            new_ids_to_consider.remove(doc_id)
            print(f"removed doc_id {doc_id} of input vector from new_ids_to_consider")

        # Loop through the subset of document vectors to consider
        for i, new_id in enumerate(new_ids_to_consider):

            # if the document ID is not yet in the set ids_to_consider...
            if new_id not in ids_to_consider_set:
                # access document_vectors_l list at index i to get the embedding
                # then append it to the list of vectors to consider as possible
                # neighbors
                document_vector_at_i = document_vectors_l[i]
                vecs_to_consider_l.append(document_vector_at_i)

            # append the new_id (the index for the document) to the list of
            # document IDs to consider
            ids_to_consider_l.append(new_id)
```

```

# also add the new_id to the set of ids to consider
# (use this to check if new_id is not already in the IDs to consider)
ids_to_consider_set.add(new_id)

### END CODE HERE ###

# Now run k-NN on the smaller set of vecs-to-consider.
print("Fast considering %d vecs" % len(vecs_to_consider_1))

# convert the vecs to consider set to a list, then to a numpy array
vecs_to_consider_arr = np.array(vecs_to_consider_1)

# call nearest neighbors on the reduced list of candidate vectors
nearest_neighbor_idx_1 = nearest_neighbor(v, vecs_to_consider_arr, k=k)

# Use the nearest neighbor index list as indices into the ids to consider
# create a list of nearest neighbors by the document ids
nearest_neighbor_ids = [ids_to_consider_1[idx]
                        for idx in nearest_neighbor_idx_1]

return nearest_neighbor_ids

```

In [42]: #document\_vecs, ind2Tweet

```

doc_id = 0
doc_to_search = all_tweets[doc_id]
vec_to_search = document_vecs[doc_id]

```

In [43]: # UNQ\_C22 (UNIQUE CELL IDENTIFIER, DO NOT EDIT)

# You do not have to input any code in this cell, but it is relevant to grading,

# Sample

```

nearest_neighbor_ids = approximate_knn(
    doc_id, vec_to_search, planes_1, k=3, num_universes_to_use=5)

```

```

removed doc_id 0 of input vector from new_ids_to_search
Fast considering 77 vecs

```

```
In [44]: print(f"Nearest neighbors for document {doc_id}")  
print(f"Document contents: {doc_to_search}")  
print("")  
  
for neighbor_id in nearest_neighbor_ids:  
    print(f"Nearest neighbor at document id {neighbor_id}")  
    print(f"document contents: {all_tweets[neighbor_id]}")
```

Nearest neighbors for document 0  
Document contents: #FollowFriday @France\_Inte @PKuchly57 @Milipol\_Paris for being top engaged members in my community this week :)

Nearest neighbor at document id 2140  
document contents: @PopsRamjet come one, every now and then is not so bad :)  
Nearest neighbor at document id 701  
document contents: With the top cutie of Bohol :) <https://t.co/Jh7F6U46UB> (<http://t.co/Jh7F6U46UB>)  
Nearest neighbor at document id 51  
document contents: #FollowFriday @France\_Espana @reglisse\_menthe @CCI\_inter for being top engaged members in my community this week :)

## 4 Conclusion

Congratulations - Now you can look up vectors that are similar to the encoding of your tweet using LSH!